

Q1 Interval scheduling and partitions revisited

Given m machines and a set of intervals S with start times s_i and finishing times f_i , output a feasible mapping $\sigma : S \rightarrow \{0, 1, 2, \dots, m\}$ where $\sigma(I) = k > 0$ means that interval I is scheduled on machine k and $\sigma(I) = 0$ means that interval I is not scheduled. The mapping is feasible if $\sigma(I_i) = \sigma(I_j) = k > 0$ implies I_i and I_j do not intersect (the intervals are compatible). The objective is to maximize $|I : \sigma(I) > 0|$.

solution

Algorithm: Sort the intervals such that $f_1 \leq f_2 \leq \dots \leq f_n$. For $i = 1 \dots n$ we determine $\sigma(I_i)$ as follows. For machine $k \in [m]$ let J_{ik} be the most recent interval that has been assigned to machine k so far, if such an interval exists. More formally, $J_{ik} = I_\ell$ for the maximum $\ell < i$ such that $\sigma(I_\ell) = k$, if such an ℓ exists. If $J_{ik} = I_\ell$ is well defined then let $a_{ik} = f_\ell$ (“ a ” for “available”), i.e. a_{ik} is the finishing time of the most recent interval assigned to machine k . Otherwise let $a_{ik} = 0$. Finally, if $a_{ik} > s_i$ for all k then let $\sigma(I_i) = 0$; otherwise let $\sigma(I_i) = k$ for some k that maximizes a_{ik} subject to $a_{ik} \leq s_i$.

Correctness: Let G_i be the greedy schedule after the first i intervals have been processed. We prove by induction on $i = 0, \dots, n$ that there exists an optimal solution OPT_i that extends G_i (in other words, OPT_i and G_i assign the first i intervals in the same way). When $i = n$ there are no more intervals to schedule so G_n must be optimal.

The base case is $i = 0$ and $G_0 = \emptyset$, so clearly there exists an optimal solution OPT_0 that extends G_0 . Assume the inductive hypothesis holds for $i - 1$. We prove it for i by considering the following cases:

- I_i is not scheduled on any machine in G_i . This implies that no legal extension of G_{i-1} can schedule I_i on any machine. In particular, OPT_{i-1} does not schedule I_i on any machine. Since OPT_{i-1} extends G_{i-1} , it follows that OPT_{i-1} extends G as well.
- I_i is scheduled on the same machine in G_i and OPT_{i-1} . Then clearly OPT_{i-1} extends G_i .
- I_i is scheduled on some machine k in G_i , but on no machine in OPT_{i-1} . Let I_j be the next interval scheduled on machine k in OPT_{i-1} , if such an interval exists.¹ Define OPT_i by replacing I_j with I_i in OPT_{i-1} . This is possible because of the following two observations:
 - At the start time s_i of I_i , machine k is free in OPT_{i-1} . This is because OPT_{i-1} extends G_{i-1} , and by the construction of G_i , at time s_i machine k is free in G_{i-1} .

¹Such an I_j must exist, because otherwise we could add I_i to OPT_{i-1} , contradicting the optimality of OPT_{i-1} .

- At the finish time f_i of I_i , machine k is free in OPT_{i-1} if I_j is removed. This is because $f_i \leq f_j$.
- I_i is scheduled on some machine k in G_i , and some different machine k' in OPT_i . Construct OPT_i as follows: start with OPT_{i-1} , and for all $j \geq i$, if I_j is scheduled on machine k (resp. k') then move it to machine k' (resp. k). This results in a valid ordering because $a_{ik'} \leq a_{ik} \leq s_i$, by the nature of the greedy algorithm

Q2 Making change

Consider the problem of making change for some amount of money n , given denominations $1 = C[1] < C[2] < \dots < C[k]$ (for example, Canadian coins come in denominations $C[1] = 1, C[2] = 5, C[3] = 10, C[4] = 25, C[5] = 100, C[6] = 200$). Say that we want the output to be a sequence of coins $S = u_1, u_2, \dots, u_m$ (where each $u_i = C[j]$ for some j) such that we use as few coins as possible (i.e., $n = u_1 + u_2 + \dots + u_m$ and m is minimal).

Give a greedy algorithm that makes change for any amount $n \geq 0$ using as few coins as possible, for the Canadian denominations. What is the runtime of your algorithm? Prove that your algorithm always returns an optimal answer.

Does your algorithm still work correctly if we add a 75 cents coin, i.e., for denominations $C[1] = 1, C[2] = 5, C[3] = 10, C[4] = 25, C[5] = 75, C[6] = 100, C[7] = 200$? What if we add a 30 cents coin instead of a 75 cents coin?

solution

Algorithm: For j from k down to 1, add $\lfloor n/C[j] \rfloor$ copies of $C[j]$ to the set of coins to be output, and update n to $(n \bmod C[j])$.

Runtime: $O(k)$, assuming that arithmetic takes constant time, and that we're allowed to express the output as an array $A[1 \dots k]$ where $A[j]$ is the number of copies of $C[j]$ in the sequence. If k is constant and we're required to write the sequence $S = u_1, u_2, \dots, u_m$ explicitly, then the runtime is $O(n)$.

Correctness: It will be more convenient to analyze the following algorithm, which is equivalent to the above but has a different runtime. For j from k down to 1, while $C[j] \leq n$, add a copy of $C[j]$ to the output and update n to $n - C[j]$.

Observe the following:

- No optimal solution has more than four 1-cent coins, because five 1-cent coins can be replaced with a 5-cent coin.
- No optimal solution has more than one 5-cent coin, because two 5-cent coins can be replaced with a 10-cent coin.

- No optimal solution has more than two (5 or 10)-cent coins, because three 10-cent coins can be replaced with a 25-cent coin and a 5-cent coin, and two 10-cent coins and a 5-cent coin can be replaced with a 25-cent coin.
- No optimal solution has more than three 25-cent coins, because four 25-cent coins can be replaced with a 100-cent coin.
- No optimal solution has more than two 100-cent coins, because two of them can be replaced with a 200-cent coin.

Now assume that this algorithm does not always produce an optimal solution. This implies that there exist n and j such that $C[j] \leq n$ and $C[j+1] > n$, but the optimal change for n does not have any copy of $C[j]$. Consider the following cases:

- $C[j] = 1$. Then $1 \leq n \leq 4$, so clearly we need 1-cent coins.
- $C[j] = 5$. Then $5 \leq n < 10$, so an optimal solution needs at least five 1-cent coins, which is a contradiction.
- $C[j] = 10$. Then $10 \leq n < 25$. But in an optimal solution, there are at most four 1-cent coins and at most one 5-cent coin, for a total of $4 * 1 + 1 * 5 = 9 < 10 \leq n$ cents, a contradiction.
- $C[j] = 25$. Then $25 \leq n < 100$. But in an optimal solution, there are at most two (5 or 10)-cent coins and at most four 1-cent coins, for a total of $4 * 1 + 2 * 10 = 24 < 25 \leq n$ cents, a contradiction.
- $C[j] = 100$. Then $100 \leq n < 200$. But in an optimal solution, there are at most three 25-cent coins, two (5 or 10)-cent coins, and four 1-cent coins, for a total of $3 * 25 + 2 * 10 + 4 * 1 = 99 < 100 \leq n$ cents, a contradiction.
- $C[j] = 200$. Then $200 \leq n$. An optimal solution has at most one 100-cent coin. By similar reasoning as above, an optimal solution has at most 199 cents, a contradiction.

If we add a 75 cents coin, then greedy does not work for $n = 150$, because the optimal solution is $75 + 75$ but greedy gives $100 + 25 + 25$.

If instead we add a 30 cents coin, then greedy does not work for $n = 50$, because the optimal solution is $25 + 25$ but greedy gives $30 + 10 + 10$.

Q3 MSTs

(a) Prove or disprove: If e is a minimum-weight edge in connected graph G (where not all edge weights are necessarily distinct), then every minimum spanning tree of G contains e .

(b) Does your answer change if $w(e)$ is unique (no other edge in G has the same weight as e , but $w(e)$ is still the smallest)? Again, prove your answer.

solution

(a) False. Let G be a cycle in which all edges have the same weight. Removing any one edge of G creates an MST.

(b) Yes, every MST contains e in this case. This follows from the correctness of Kruskal's algorithm. For an alternate proof: let T be a spanning tree without e . Adding e to T creates a cycle in T , and removing any other edge from that cycle creates a spanning tree with lower weight.

Q4 Cops and robbers

Given an array of size n that has the following specifications: Each element in the array contains either a cop or a robber. Each cop can catch only one robber. A cop cannot catch a robber who is more than K units away from the cop. Write an algorithm to find the maximum number of robbers that can be caught.

solution

Algorithm: Initialize c to the first cop in the array, and r to the first robber in the array. While c and r are not null, do the following. If $|c - r| \leq K$ then assign c to catch r , and increment c and r to the next cop and robber respectively. Else, if $c < r$ (resp. $r < c$) then increment c to the next cop (resp. r to the next robber).

Correctness: Let G_i be the greedy assignment after the i 'th iteration of the greedy algorithm. We prove by induction on i that there exists an optimal assignment OPT_i that extends G_i . When i is the last iteration there are no more robbers to assign or no more cops to assign, so this solution must be optimal.

The base case is $i = 0$ and $G_0 = \emptyset$, so clearly there exists an optimal solution OPT_0 that extends G_0 . Assume the inductive hypothesis holds for $i - 1$. Let c and r be the cop and robber considered in the i 'th iteration. We prove the inductive hypothesis for i by considering the following cases:

- G_i does not match r with c . Then $|r - c| > K$, so OPT_{i-1} doesn't match r with c either. Since OPT_{i-1} extends G_{i-1} , it follows that OPT_{i-1} extends G_i as well, so we can let $OPT_i = OPT_{i-1}$.
- G_i matches r with c , and
 - OPT_{i-1} doesn't match r , or OPT_{i-1} doesn't match c . Form OPT_i from OPT_{i-1} by unmatching r or c , and then matching r with c . This doesn't decrease the total number of matchings, so OPT_i must be optimal because OPT_{i-1} is optimal. Clearly OPT_i extends G_i .
 - OPT_{i-1} matches c with r . Then let $OPT_i = OPT_{i-1}$.
 - OPT_{i-1} matches c with r' , and r with c' , where $r \neq r'$ and $c \neq c'$. Since $r < r'$ and $c < c'$, the largest-indexed of $\{r, r', c, c'\}$ must be either r' or c' . Assume without loss of generality that it's c' . Then since OPT_{i-1} matches c' and r , it follows that

$K \geq |c' - r| = c' - r > c' - r' = |c' - r'|$. Therefore we can define OPT_i by starting with OPT_{i-1} , and replacing the matchings (r, c') and (r', c) with (r, c) and (r', c') .